

# HateClassification

June 10, 2023

```
[1]: import numpy as np
import pandas as pd
import string
import re
import nltk
nltk.download('punkt')
nltk.download('stopwords')
nltk.download('wordnet')
```

```
[nltk_data] Downloading package punkt to /home/sit/nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package stopwords to /home/sit/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to /home/sit/nltk_data...
[nltk_data] Package wordnet is already up-to-date!
```

[1]: True

```
[2]: train_data = pd.read_excel('constraint_Hindi_Train.xlsx')
valid_data = pd.read_excel('Constraint_Hindi_Valid.xlsx')
train_data.head()
```

```
[2]:
```

	Unique ID	Post	\
0	1	...	
1	2	...	
2	3	9 , ...	
3	4	@prabhav218 ...	
4	5	#unlock4guidelines - -4 ...	

  

	Labels Set
0	hate,offensive
1	non-hostile
2	non-hostile
3	defamation,offensive
4	non-hostile

```
[3]: train_data.rename(columns={'Unique ID':'ID','Post':'Text','Labels Set':
    ↳ 'Labels'},inplace=True)
```

```
valid_data.rename(columns={'Unique ID':'ID','Post':'Text','Labels Set':
↳ 'Labels'},inplace=True)
train_data.head()
```

```
[3]:   ID                               Text          Labels
0    1                ...      hate,offensive
1    2                ...      non-hostile
2    3          9      , ...      non-hostile
3    4  @prabhav218      ...  defamation,offensive
4    5  #unlock4guidelines - -4      ...      non-hostile
```

```
[4]: print(train_data.Labels.nunique())
print(train_data.Labels.unique())
print(train_data.Labels.value_counts())
```

```
16
['hate,offensive' 'non-hostile' 'defamation,offensive' 'fake' 'hate'
 'offensive' 'fake,hate' 'defamation' 'defamation,hate'
 'defamation,hate,offensive' 'defamation,fake,offensive' 'fake,offensive'
 'defamation,fake' 'defamation,fake,hate' 'fake,hate,offensive'
 'defamation,fake,hate,offensive']
non-hostile                3050
fake                        1009
hate                        478
offensive                   405
defamation                  305
hate,offensive              163
defamation,offensive         81
defamation,hate              74
defamation,fake              34
defamation,hate,offensive    28
fake,offensive               28
fake,hate                    27
defamation,fake,offensive    24
defamation,fake,hate         9
defamation,fake,hate,offensive 9
fake,hate,offensive          4
Name: Labels, dtype: int64
```

```
[5]: train_sample = train_data.loc[train_data['Labels'].
↳ isin(['non-hostile','fake','hate','offensive','defamation'])]
valid_sample = valid_data.loc[valid_data['Labels'].
↳ isin(['non-hostile','fake','hate','offensive','defamation'])]
train_sample.head()
```

```
[5]:   ID                               Text          Labels
1    2                ...      non-hostile
2    3          9      , ...      non-hostile
```

```

4 5 #unlock4guidelines - -4 ... non-hostile
5 6 UN ... fake
6 7 #Corona \n#ZeeJankar... non-hostile

```

```

[6]: print(len(train_sample))
      print(len(valid_sample))

```

5247

747

```

[7]: import re
      import string
      import nltk
      from nltk.corpus import stopwords
      from nltk.tokenize import word_tokenize
      from nltk.stem import WordNetLemmatizer
      import emoji
      #from cleantext import clean

      def preprocess_tweet(tweet):
          # Convert to lowercase
          tweet = tweet.lower()

          # Remove URLs
          tweet = re.sub(r"http\S+|www\S+|https\S+", "", tweet, flags=re.MULTILINE)

          # Remove usernames and hashtags symbols
          tweet = re.sub(r"@w+|#+", "", tweet)

          #Clean tweet
          #tweet = clean(tweet,no_emoji=True)

          # Remove punctuation
          tweet = tweet.translate(str.maketrans("", "", string.punctuation))

          #Encoding UTF-8
          #tweet = tweet.encode(encoding='UTF-8',errors='strict')
          # Tokenize the tweet
          #tokens = word_tokenize(tweet)

          # Remove stopwords
          #stop_words = set(stopwords.words("english"))
          #tokens = [token for token in tokens if token not in stop_words]

          # Lemmatize tokens
          #lemmatizer = WordNetLemmatizer()
          #tokens = [lemmatizer.lemmatize(token) for token in tokens]

```

```

    # Join tokens back into a single string
    processed_tweet = "".join(str(tweet))

    return processed_tweet

# Example usage
tweet = "Great article on #AI! @username Check it out: https://www.example.com/
→article"
processed_tweet = preprocess_tweet(tweet)
print(processed_tweet)

```

great article on ai check it out

```

[8]: train_sample['processed_text'] = train_sample['Text'].apply(lambda x :
    →preprocess_tweet(x))
train_sample.head()

```

/tmp/ipykernel\_5136/21834224.py:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
[docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```

train_sample['processed_text'] = train_sample['Text'].apply(lambda x :
preprocess_tweet(x))

```

```

[8]:      ID                               Text      Labels \
1      2      ... non-hostile
2      3      9      , ... non-hostile
4      5  #unlock4guidelines - -4      ... non-hostile
5      6      UN      ...      fake
6      7      #Corona      \n#ZeeJankar... non-hostile

                                processed_text
1      ...
2      9      ...
4  unlock4guidelines  4      ...
5      un      ...
6      corona      \nzeejankario...

```

```

[9]: valid_sample['processed_text'] = valid_sample['Text'].apply(lambda x :
    →preprocess_tweet(x))
valid_sample.head()

```

/tmp/ipykernel\_5136/3194795115.py:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: <https://pandas.pydata.org/pandas->

```
docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
valid_sample['processed_text'] = valid_sample['Text'].apply(lambda x :
preprocess_tweet(x))
```

```
[9]: ID                                Text      Labels \
0    1                ... non-hostile
1    2      rss          ... defamation
2    3      /   10 ... non-hostile
3    4    PM - ... non-hostile
4    5 :   Toilet ,    ... non-hostile

                                processed_text
0                ...
1      rss          ...
2      10 ...
3    pm          ...
4    toilet      ...
```

```
[10]: #Unique Word Count
from collections import Counter
frames = [train_sample,valid_sample]
words_df = pd.concat(frames,axis=0)

results = Counter()
words_df['processed_text'].str.split().apply(results.update)
print(len(results))
```

20916

```
[11]: from sklearn.feature_extraction.text import TfidfVectorizer
tfidf = TfidfVectorizer()
```

```
[12]: from sklearn.preprocessing import LabelEncoder
lbl_enc = LabelEncoder()
```

```
[13]: X_train_vec = tfidf.fit_transform(train_sample['processed_text'])
X_valid_vec = tfidf.transform(valid_sample['processed_text'])
```

```
[14]: lbl_enc.fit(train_sample['Labels'])
train_sample['Enc_Labels'] = lbl_enc.transform(train_sample['Labels'])
valid_sample['Enc_Labels'] = lbl_enc.transform(valid_sample['Labels'])
```

/tmp/ipykernel\_5136/2798184267.py:2: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
train\_sample['Enc\_Labels'] = lbl\_enc.transform(train\_sample['Labels'])

/tmp/ipykernel\_5136/2798184267.py:3: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
valid\_sample['Enc\_Labels'] = lbl\_enc.transform(valid\_sample['Labels'])

```
[15]: lbl_enc.classes_
```

```
[15]: array(['defamation', 'fake', 'hate', 'non-hostile', 'offensive'],  
       dtype=object)
```

```
[16]: from sklearn.ensemble import RandomForestClassifier  
rfc = RandomForestClassifier()  
rfc.fit(X_train_vec, train_sample['Enc_Labels'])
```

```
[16]: RandomForestClassifier()
```

```
[17]: y_pred = rfc.predict(X_valid_vec)
```

```
[18]: from sklearn.metrics import classification_report  
print(classification_report(valid_sample['Enc_Labels'], y_pred))
```

	precision	recall	f1-score	support
0	0.33	0.02	0.04	43
1	0.52	0.33	0.40	144
2	0.22	0.06	0.09	68
3	0.68	0.95	0.79	435
4	0.35	0.19	0.25	57
accuracy			0.64	747
macro avg	0.42	0.31	0.32	747
weighted avg	0.56	0.64	0.57	747

```
[41]: import matplotlib.pyplot as plt  
import os  
import re  
import shutil  
import string  
import tensorflow as tf  
  
from tensorflow.keras.preprocessing.text import Tokenizer  
from tensorflow.keras.preprocessing.sequence import pad_sequences  
from tensorflow.keras import layers  
from tensorflow.keras import losses
```

```
[42]: train_text = train_sample['processed_text'].values
train_label = train_sample['Enc_Labels'].values
valid_text = valid_sample['processed_text'].values
valid_label = valid_sample['Enc_Labels'].values

[43]: # Define the maximum number of words to keep in the vocabulary
max_words = 10000

# Create a tokenizer and fit it on the training text data
tokenizer = Tokenizer(num_words=max_words)
tokenizer.fit_on_texts(train_text)

# Convert the text data to sequences
train_sequences = tokenizer.texts_to_sequences(train_text)
valid_sequences = tokenizer.texts_to_sequences(valid_text)

# Pad the sequences to have the same length
max_sequence_length = max([len(sequence) for sequence in train_sequences])
train_data = pad_sequences(train_sequences, maxlen=max_sequence_length)
valid_data = pad_sequences(valid_sequences, maxlen=max_sequence_length)

[44]: # Convert the labels to one-hot encoded vectors
num_classes = len(np.unique(train_label))
train_label = tf.keras.utils.to_categorical(train_label, num_classes=num_classes)
valid_label = tf.keras.utils.to_categorical(valid_label, num_classes=num_classes)

[45]: # Define the model architecture
model = tf.keras.models.Sequential([
    tf.keras.layers.Embedding(input_dim=max_words, output_dim=100,
    ↪input_length=max_sequence_length),
    tf.keras.layers.Conv1D(128, 5, activation='relu'),
    tf.keras.layers.GlobalMaxPooling1D(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(num_classes, activation='softmax')
])

[55]: model.compile(loss='categorical_crossentropy', optimizer='adam',
    ↪metrics=['accuracy', 'categorical_accuracy'])

[56]: # Train the model
model.fit(train_data, train_label, epochs=10, batch_size=32,
    ↪validation_data=(valid_data, valid_label))
```

```
Epoch 1/10
164/164 [=====] - 6s 30ms/step - loss: 0.0146 -
accuracy: 0.9975 - categorical_accuracy: 0.9975 - val_loss: 1.5407 -
val_accuracy: 0.7122 - val_categorical_accuracy: 0.7122
Epoch 2/10
```

```

164/164 [=====] - 5s 29ms/step - loss: 0.0134 -
accuracy: 0.9966 - categorical_accuracy: 0.9966 - val_loss: 1.5798 -
val_accuracy: 0.7082 - val_categorical_accuracy: 0.7082
Epoch 3/10
164/164 [=====] - 5s 29ms/step - loss: 0.0089 -
accuracy: 0.9977 - categorical_accuracy: 0.9977 - val_loss: 1.5667 -
val_accuracy: 0.7149 - val_categorical_accuracy: 0.7149
Epoch 4/10
164/164 [=====] - 5s 29ms/step - loss: 0.0077 -
accuracy: 0.9975 - categorical_accuracy: 0.9975 - val_loss: 1.7107 -
val_accuracy: 0.7149 - val_categorical_accuracy: 0.7149
Epoch 5/10
164/164 [=====] - 5s 29ms/step - loss: 0.0070 -
accuracy: 0.9981 - categorical_accuracy: 0.9981 - val_loss: 1.7269 -
val_accuracy: 0.7055 - val_categorical_accuracy: 0.7055
Epoch 6/10
164/164 [=====] - 5s 29ms/step - loss: 0.0087 -
accuracy: 0.9979 - categorical_accuracy: 0.9979 - val_loss: 1.8557 -
val_accuracy: 0.7162 - val_categorical_accuracy: 0.7162
Epoch 7/10
164/164 [=====] - 5s 29ms/step - loss: 0.0059 -
accuracy: 0.9977 - categorical_accuracy: 0.9977 - val_loss: 2.1018 -
val_accuracy: 0.7216 - val_categorical_accuracy: 0.7216
Epoch 8/10
164/164 [=====] - 5s 29ms/step - loss: 0.0083 -
accuracy: 0.9970 - categorical_accuracy: 0.9970 - val_loss: 1.9744 -
val_accuracy: 0.6908 - val_categorical_accuracy: 0.6908
Epoch 9/10
164/164 [=====] - 5s 29ms/step - loss: 0.0046 -
accuracy: 0.9985 - categorical_accuracy: 0.9985 - val_loss: 1.9756 -
val_accuracy: 0.7149 - val_categorical_accuracy: 0.7149
Epoch 10/10
164/164 [=====] - 5s 29ms/step - loss: 0.0036 -
accuracy: 0.9985 - categorical_accuracy: 0.9985 - val_loss: 1.9839 -
val_accuracy: 0.7175 - val_categorical_accuracy: 0.7175

```

[56]: <keras.callbacks.History at 0x7f6aa8559b50>

```

[48]: # Evaluate the model on the test set
loss, accuracy = model.evaluate(valid_data, valid_label)
print(f'Test loss: {loss:.4f}')
print(f'Test accuracy: {accuracy:.4f}')

```

```

24/24 [=====] - 0s 11ms/step - loss: 1.1967 - accuracy:
0.7202
Test loss: 1.1967
Test accuracy: 0.7202

```



```
[57]: pred = model.predict(valid_data)
```

```
24/24 [=====] - 0s 11ms/step
```

```
[60]: print(classification_report(np.argmax(valid_label,axis=1),np.  
    ↪argmax(pred,axis=1)))
```

	precision	recall	f1-score	support
0	0.44	0.33	0.37	43
1	0.52	0.63	0.57	144
2	0.43	0.38	0.40	68
3	0.88	0.88	0.88	435
4	0.52	0.40	0.46	57
accuracy			0.72	747
macro avg	0.56	0.52	0.54	747
weighted avg	0.72	0.72	0.71	747

```
[ ]:
```